

**LDRP Institute of technology and Research**  
**Computer Engineering Department**  
**Subject Name: Natural Language Processing(CT703C-N)**

**Practical List**

<b>Sr. No</b>	<b>Practical</b>	<b>Date</b>	<b>Pg. No.</b>	<b>Sign</b>
1	Basic Text Processing operation on text document.	13/07/2022		
2	Implement N-gram Language model.	20/07/2022		
3	Write a program to extract features from text.	27/07/2022		
4	Implement word embedding using Word2Vec/Glove/FasText	03/08/2022		
5	Implement LSA and Topic model.	10/08/2022		
6	Implementation text classification using Naive Bayes, SVM.	17/08/2022		
7	Implementation of K-means Clustering algorithm on text.	24/08/2022		
8	Implement PoS Tagging on text.	07/09/2022		
9	Implement text processing with neural network.	14/09/2022		
10	Implement text processing with LSTM.	21/09/2022		
11	Implement HMM/CRF on sequence tagging task.	28/09/2022		

## **PRACTICAL - 1**

**Aim:** Basic Text Processing operation on text document.

### **Code:**

```
import nltk

nltk.download('punkt')

text = "Backgammon is one of the oldest known board games. Its history can be traced back
nearly 5,000 years to archeological discoveries in the Middle East. It is a two player game
where each player has fifteen checkers which move between twenty-four points according
to the roll of two dice."

# Sentence tokenization

sentences = nltk.sent_tokenize(text)

print("Sentence Tokenization :")

for sentence in sentences:
    print(sentence)
    print()

# Word tokenization

print("Word Tokenization :")

for sentence in sentences:
    words = nltk.word_tokenize(sentence)
    print(words)
    print()

# Stemming and Lemmatization

from nltk.stem import PorterStemmer, WordNetLemmatizer

from nltk.corpus import wordnet

nltk.download('wordnet')

nltk.download('omw-1.4')

def compare_stemmer_and_lemmatizer(stemmer, lemmatizer, word, pos):
```

```

    print("Stemmer:", stemmer.stem(word))

    print("Lemmatizer:", lemmatizer.lemmatize(word, pos))

    print()

    lemmatizer = WordNetLemmatizer()

    stemmer = PorterStemmer()

    compare_stemmer_and_lemmatizer(stemmer, lemmatizer, word = "seen", pos =
wordnet.VERB)

    compare_stemmer_and_lemmatizer(stemmer, lemmatizer, word = "drove", pos =
wordnet.VERB)


# Stopword Removal

from nltk.corpus import stopwords

nltk.download('stopwords')

stop_words = set(stopwords.words("english"))

sentence = "Backgammon is one of the oldest known board games."

print("Stop word removal")

words = nltk.word_tokenize(sentence)

without_stop_words = [word for word in words if not word in stop_words]

print(without_stop_words)


# Regex

import re

sentence = "The development of snowboarding was inspired by skateboarding, sledding,
surfing and skiing."

pattern = r"^[^w]"

print("Regex")

print(re.sub(pattern, " ", sentence))


# Bag of Word

from sklearn.feature_extraction.text import CountVectorizer

import pandas as pd

documents = ["I like this movie, it's funny.", 'I hate this movie.', 'This was awesome! I like it.',
'Nice one. I love it.']

```

```
count_vectorizer = CountVectorizer()

bag_of_words = count_vectorizer.fit_transform(documents)

feature_name = count_vectorizer.get_feature_names()

print("Bag-of-Words")

pd.DataFrame(bag_of_words.toarray(), columns = feature_name)


# TF-IDF

from sklearn.feature_extraction.text import TfidfVectorizer

import pandas as pd

tfidf_vectorizer = TfidfVectorizer()

values = tfidf_vectorizer.fit_transform(documents)

feature_names = tfidf_vectorizer.get_feature_names()

print("TF-IDF")

pd.DataFrame(values.toarray(), columns = feature_names)
```

## Output:

```
Sentence Tokenization :
Backgammon is one of the oldest known board games.

Its history can be traced back nearly 5,000 years to archeological discoveries in the Middle East.

Word Tokenization :
['Backgammon', 'is', 'one', 'of', 'the', 'oldest', 'known', 'board', 'games', '.']

['Its', 'history', 'can', 'be', 'traced', 'back', 'nearly', '5,000', 'years', 'to', 'archeological', 'discoveries', 'in', 'the', 'Middle', 'East', '.']
```

```
Stemmer: seen
Lemmatizer: see

Stemmer: drove
Lemmatizer: drive
```

```
Stop word removal
['Backgammon', 'one', 'oldest', 'known', 'board', 'games', '.']
```

```
Regex
The development of snowboarding was inspired by skateboarding  sledding  surfing and skiing
```

	awesome	funny	hate	it	like	love	movie	nice	one	this	was
0	0	1	0	1	1	0	1	0	0	1	0
1	0	0	1	0	0	0	1	0	0	1	0
2	1	0	0	1	1	0	0	0	0	1	1
3	0	0	0	1	0	1	0	1	1	0	0

	awesome	funny	hate	it	like	love	movie	nice	one	this	was
0	0.000000	0.571848	0.000000	0.365003	0.450852	0.000000	0.450852	0.000000	0.000000	0.365003	0.000000
1	0.000000	0.000000	0.702035	0.000000	0.000000	0.000000	0.553492	0.000000	0.000000	0.448100	0.000000
2	0.539445	0.000000	0.000000	0.344321	0.425305	0.000000	0.000000	0.000000	0.000000	0.344321	0.539445
3	0.000000	0.000000	0.000000	0.345783	0.000000	0.541736	0.000000	0.541736	0.541736	0.000000	0.000000

## **PRACTICAL - 2**

**Aim:** Implement N-gram Language model.

### **Code:**

```
from nltk.corpus import reuters
from nltk import bigrams, trigrams
from collections import Counter, defaultdict

model = defaultdict(lambda: defaultdict(lambda: 0))

for sentence in reuters.sents():
    for w1, w2, w3 in trigrams(sentence, pad_right=True, pad_left=True):
        model[(w1, w2)][w3] += 1

for w1_w2 in model:
    total_count = float(sum(model[w1_w2].values()))
    for w3 in model[w1_w2]:
        model[w1_w2][w3] /= total_count

import random

text = ["today", "the"]
sentence_finished = False

while not sentence_finished:
    r = random.random()
    accumulator = .0

    for word in model[tuple(text[-2:]).keys():
        accumulator += model[tuple(text[-2:])[word]
```

```
        if accumulator >= r:
            text.append(word)
            break

    if text[-2:] == [None, None]:
        sentence_finished = True

    print(' '.join([t for t in text if t]))
```

**Output:**

```
today the price to rise in import duties on frozen orange juice imports have ratified it .
```

## **PRACTICAL - 3**

**Aim:** Write a program to extract features from text.

### **Code:**

```
A = [' Messi is running towards the Goalpost #football.',
'Ronaldo is better than Messi',
'Messi is better than Ronaldo',
'Messi is the no 1 football player',
'Messi Messi Ronaldo Ronaldo',
'mbappe,mbappe,mbappe,mbappe,mbappe,mbappe,mbappe,mbappe,mbappe,']

import pandas as pd
data = pd.DataFrame({'tweet_text':A})

data['word_count'] = data['tweet_text'].apply(lambda x: len(str(x).split(" ")))
data['char_count'] = data['tweet_text'].str.len() ## this also includes spaces

def avg_word(sentence):
    words = sentence.split()
    return (sum(len(word) for word in words)/len(words))
data['avg_word'] = data['tweet_text'].apply(lambda x: avg_word(x))

import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
stop = stopwords.words('english')

data['stopwords'] = data['tweet_text'].apply(lambda x: len([x for x in x.split() if x in stop]))
data['hashtags'] = data['tweet_text'].apply(lambda x: len([x for x in x.split() if x.startswith('#')]))
data['numerics'] = data['tweet_text'].apply(lambda x: len([x for x in x.split() if x.isdigit()]))
data['uppercase character'] = data['tweet_text'].apply(lambda x: len([x for x in x.split() if x.isupper()])))
```



```
pos_family = {
    'noun': ['NN','NNS','NNP','NNPS'],
    'pron': ['PRP','PRP$','WP','WP$'],
    'verb': ['VB','VBD','VBG','VBN','VBP','VBZ'],
    'adj': ['JJ','JJR','JJS'],
    'adv': ['RB','RBR','RBS','WRB']
}

from textblob import TextBlob, Word, Blobber
import nltk
nltk.download('punkt')
nltk.download('averaged_perceptron_tagger')

def check_pos_tag(x, flag):
    cnt = 0
    try:
        wiki = TextBlob(x)
        for tup in wiki.tags:
            ppo = list(tup)[1]
            if ppo in pos_family[flag]:
                cnt += 1
    except:
        pass
    return cnt

data['noun_count'] = data['tweet_text'].apply(lambda x: check_pos_tag(x, 'noun'))
data['verb_count'] = data['tweet_text'].apply(lambda x: check_pos_tag(x, 'verb'))
data['adj_count'] = data['tweet_text'].apply(lambda x: check_pos_tag(x, 'adj'))
data['adv_count'] = data['tweet_text'].apply(lambda x: check_pos_tag(x, 'adv'))
data['pron_count'] = data['tweet_text'].apply(lambda x: check_pos_tag(x, 'pron'))

data.head()
```

**Output:**

	tweet_text	word_count	char_count	avg_word	stopwords	hashtags	numerics	uppercase character	noun_count	verb_count	adj_count	adv_count	pron_count
0	Messi is running towards the Goalpost #football.	8	49	6.000000	2	1	0	0	3	2	0	0	0
1	Ronaldo is better than Messi	5	28	4.800000	2	0	0	0	2	1	1	0	0
2	Messi is better than Ronaldo	5	28	4.800000	2	0	0	0	2	1	1	0	0
3	Messi is the no 1 football player	7	33	3.857143	3	0	1	0	3	1	0	0	0
4	Messi Messi Ronaldo Ronaldo	4	27	6.000000	0	0	0	0	4	0	0	0	0

## **PRACTICAL - 4**

**Aim:** Implement word embedding using Word2Vec/Glove/fastText.

**Code:**

```
from gensim.models import Word2Vec

sentences = [['this', 'is', 'the', 'good', 'machine', 'learning', 'book'],
              ['this', 'is', 'another', 'machine', 'learning', 'book'],
              ['one', 'more', 'new', 'book'],
              ['this', 'is', 'about', 'machine', 'learning', 'post'],
              ['orange', 'juice', 'is', 'the', 'liquid', 'extract', 'of', 'fruit'],
              ['orange', 'juice', 'comes', 'in', 'several', 'different', 'varieties'],
              ['this', 'is', 'the', 'last', 'machine', 'learning', 'book'],
              ['orange', 'juice', 'comes', 'in', 'several', 'different', 'packages'],
              ['orange', 'juice', 'is', 'liquid', 'extract', 'from', 'fruit', 'on', 'orange', 'tree']]

from gensim.models import FastText

model = Word2Vec(sentences, size=20, min_count=1, window=2, sg=0)

is_model = model['is']
print("Model of is \n", is_model)

orange_juice = model.similarity('orange', 'juice')
print("Similarity between orange and juice is ", orange_juice)

this_orange = model.similarity('this', 'orange')
print("Similarity between this and orange is ", this_orange)

most_similar_orange = model.most_similar('orange')[:2]
print("2 most similar word to orange", most_similar_orange)
```

```

close_words = model.similar_by_word('orange')
print("Close words to orange is \n",close_words)

import numpy as np
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt

def display_closestwords_tsnescaatterplot(model, word, size):

    arr = np.empty((0,size), dtype='f')
    word_labels = [word]

    close_words = model.similar_by_word(word)
    arr = np.append(arr, np.array([model[word]]), axis=0)
    for wrd_score in close_words:
        wrd_vector = model[wrd_score[0]]
        word_labels.append(wrd_score[0])
        arr = np.append(arr, np.array([wrd_vector]), axis=0)

    tsne = TSNE(n_components=2, random_state=0)
    np.set_printoptions(suppress=True)
    Y = tsne.fit_transform(arr)
    x_coords = Y[:, 0]
    y_coords = Y[:, 1]
    plt.scatter(x_coords, y_coords)
    for label, x, y in zip(word_labels, x_coords, y_coords):
        plt.annotate(label, xy=(x, y), xytext=(0, 0), textcoords='offset points')
        plt.xlim(x_coords.min()+0.00005, x_coords.max()+0.00005)
        plt.ylim(y_coords.min()+0.00005, y_coords.max()+0.00005)
    plt.show()

display_closestwords_tsnescaatterplot(model, 'orange', 4)



```




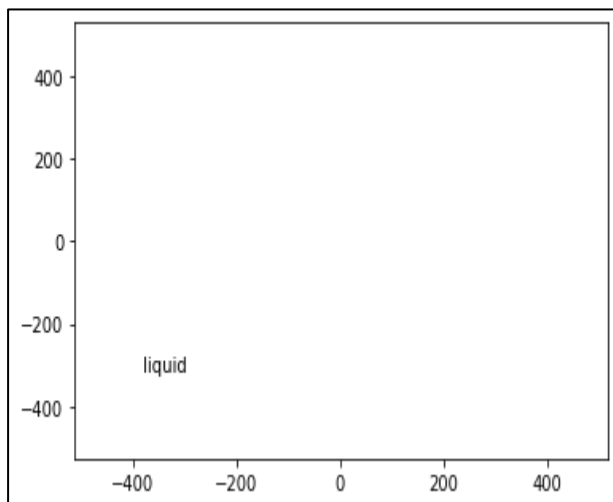


## Output:




```

Model of is
[ 0.02064108  0.02450296 -0.0184775  -0.01977053 -0.01373223  0.02217131
 -0.00491822 -0.01642907  0.01152982  0.01568795 -0.02030085 -0.01368141
 -0.0166703  -0.01665059 -0.01207751 -0.00259845 -0.01041098 -0.01049821
  0.01156926 -0.00734999]
Similarity between orange and juice is  0.2692408
Similarity between this and orange is  0.20401147
2 most similar word to orange [('in', 0.32315507531166077), ('juice', 0.2692407965660095)]
Close words to orange is
[('in', 0.32315507531166077),
 ('juice', 0.2692407965660095),
 ('varieties', 0.24400851130485535),
 ('good', 0.23437678813934326),
 ('this', 0.2040114849805832),
 ('fruit', 0.16190184652805328),
 ('learning', 0.15402346849441528),
 ('last', 0.14784149825572968),
 ('different', 0.0635468065738678),
 ('new', 0.04215109348297119)]

```







## **PRACTICAL - 5**

**Aim:** Implement LSA and Topic model.

**Code:**

```
import pandas as pd

from sklearn.feature_extraction.text import TfidfVectorizer

import nltk

from sklearn.decomposition import TruncatedSVD

nltk.download('stopwords')

from nltk.corpus import stopwords

a1 = "He is a good dog."
a2 = "The dog is too lazy."
a3 = "That is a brown cat."
a4 = "The cat is very active."
a5 = "I have brown cat and dog."

df = pd.DataFrame()

df["documents"] = [a1,a2,a3,a4,a5]

df['clean_documents'] = df['documents'].str.replace("[^a-zA-Z#]", " ")

df['clean_documents'] = df['clean_documents'].fillna("").apply(lambda x: ' '.join([w for w in
x.split() if len(w)>2]))

df['clean_documents'] = df['clean_documents'].fillna("").apply(lambda x: x.lower())

tokenized_doc = df['clean_documents'].fillna("").apply(lambda x: x.split())

tokenized_doc = tokenized_doc.apply(lambda x: [itemstop_words =
stopwords.words('english') for item in x if item not in stop_words])

detokenized_doc = []

for i in range(len(df)):

    t = ' '.join(tokenized_doc[i])

    detokenized_doc.append(t)
```



```

df['clean_documents'] = detokenized_doc

vectorizer = TfidfVectorizer(stop_words='english', smooth_idf=True)

X = vectorizer.fit_transform(df['clean_documents'])

svd_model = TruncatedSVD(n_components=2, algorithm='randomized', n_iter=100,
random_state=122)

lsa = svd_model.fit_transform(X)

pd.options.display.float_format = '{:,.16f}'.format

topic_encoded_df = pd.DataFrame(lsa, columns = ["topic_1", "topic_2"])

topic_encoded_df["documents"] = df['clean_documents']

dictionary = vectorizer.get_feature_names()

encoding_matrix = pd.DataFrame(svd_model.components_, index = ["topic_1","topic_2"],
columns = (dictionary)).T

display(topic_encoded_df[["documents", "topic_1", "topic_2"]])

display(encoding_matrix)

```

## Output:

	documents	topic_1	topic_2
0	good dog	0.3413834191239963	0.7199781067501041
1	the dog too lazy	0.3413834191239966	0.7199781067501029
2	that brown cat	0.8609490919302167	-0.3659836550739514
3	the cat very active	0.5166658991993207	-0.3850046207843261
4	have brown cat and dog	0.9494117370834869	0.0236302940661148
	topic_1	topic_2	
<b>active</b>	0.2003541259081108	-0.2424408501618362	
<b>brown</b>	0.5965117122287049	-0.2018098984872580	
<b>cat</b>	0.6293380994160952	-0.3298859088715316	
<b>dog</b>	0.4158307960649448	0.6169033286639758	
<b>good</b>	0.1323826028466491	0.4533766476433699	
<b>lazy</b>	0.1323826028466496	0.4533766476433685	

## **PRACTICAL - 6**

**Aim:** Implementation text classification using Naïve Bayes, SVM.

**Code:**

```
import numpy as np, pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.datasets import fetch_20newsgroups

from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.naive_bayes import MultinomialNB

from sklearn.svm import SVC

from sklearn.pipeline import make_pipeline

from sklearn.metrics import confusion_matrix, accuracy_score

sns.set()


data = fetch_20newsgroups()

text_categories = data.target_names

train_data = fetch_20newsgroups(subset="train", categories=text_categories)

test_data = fetch_20newsgroups(subset="test", categories=text_categories)


model = make_pipeline(TfidfVectorizer(), SVC())

model.fit(train_data.data, train_data.target)

predicted_categories = model.predict(test_data.data)


mat = confusion_matrix(test_data.target, predicted_categories)

sns.heatmap(mat.T, square = True, annot=True, fmt = "d",
xticklabels=train_data.target_names,yticklabels=train_data.target_names)

plt.xlabel("true labels")

plt.ylabel("predicted label")

plt.show()

print("The accuracy of SVM is {}".format(accuracy_score(test_data.target,
predicted_categories)))
```

```

model1 = make_pipeline(TfidfVectorizer(), MultinomialNB())

model1.fit(train_data.data, train_data.target)

predicted_categories1 = model1.predict(test_data.data)

print("Naive Bayes Confuion Matrix \n")

mat = confusion_matrix(test_data.target, predicted_categories1)

sns.heatmap(mat.T, square = True, annot=True, fmt = "d",
xticklabels=train_data.target_names,yticklabels=train_data.target_names)

plt.xlabel("true labels")

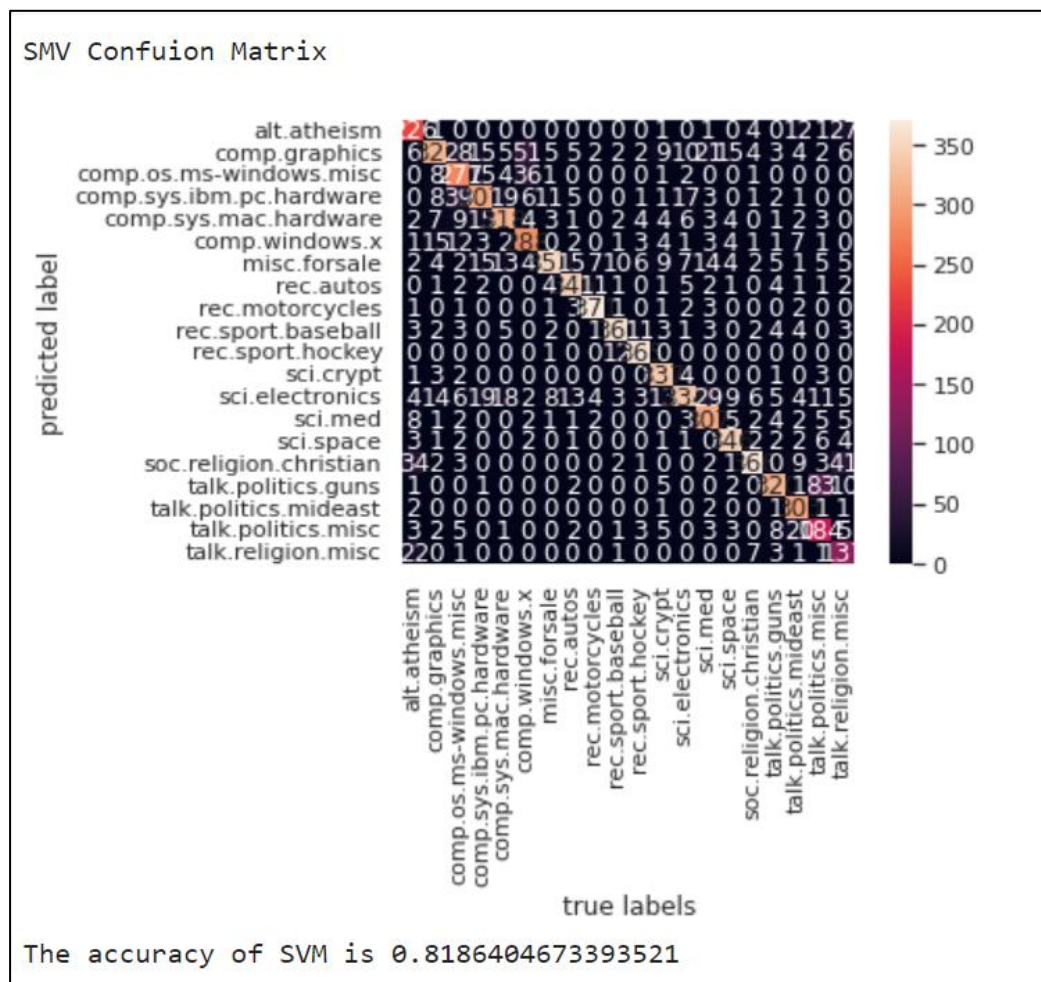
plt.ylabel("predicted label")

plt.show()

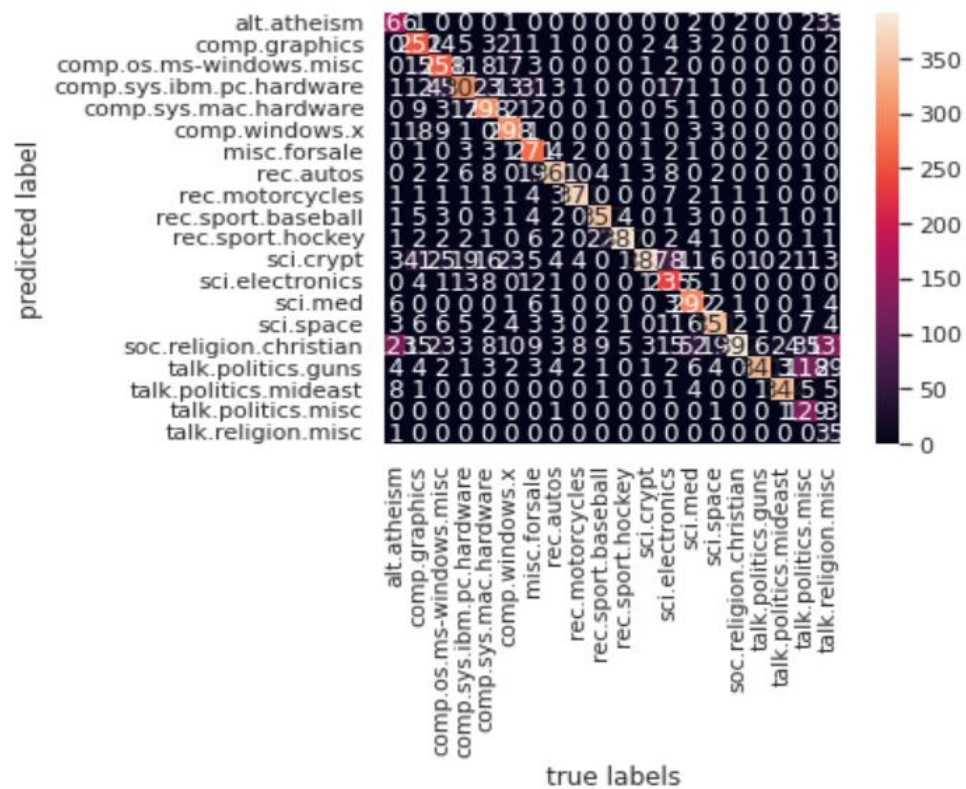
print("The accuracy of Naive Bayes is {}".format(accuracy_score(test_data.target,
predicted_categories1)))

```

## Output:



Naive Bayes Confuion Matrix



## PRACTICAL-7

**Aim:** Implementation of K-means Clustering algorithm on text.

**Code:**

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.cluster import KMeans

documents = ["the young french men crowned world champions",
             "Google Translate app is getting more intelligent everyday",
             "Facebook face recognition is driving me crazy",
             "who is going to win the Golden Ball title this year",
             "these camera apps are funny",
             "Croacian team made a brilliant world cup campaign reaching the final match",
             "Google Chrome extensions are useful.",
             "Social Media apps leveraging AI incredibly",
             "Qatar 2022 FIFA world cup is played in winter"]

vectorizer = TfidfVectorizer(stop_words = 'english')
data = vectorizer.fit_transform(documents)

clustering_model = KMeans(n_clusters = 2, init = 'k-means++', max_iter = 300, n_init = 10)
clustering_model.fit(data)

print("Top terms per cluster:")
sorted_centroids = clustering_model.cluster_centers_.argsort()[:, :-1]
terms = vectorizer.get_feature_names()

for i in range(true_k):
    print("Cluster %d:" % i, end="")
    for ind in sorted_centroids[i, :10]:
```

```
        print(' %s' % terms[ind], end = ' ')

    print()

    print()

    print()

    print("Predictions of new documents")

    new_doc1 = ["how to install Chrome"]

    Y = vectorizer.transform(new_doc1)

    prediction1 = clustering_model.predict(Y)

    print("Cluster of doc1 is ",prediction1)


    new_doc2 = ["UCL Final match is played in Madrid this year"]

    Y = vectorizer.transform(new_doc2)

    prediction2 = clustering_model.predict(Y)

    print("Cluster of doc2 is ",prediction2)
```

## Output:

```
Top terms per cluster:
Cluster 0: apps google funny camera extensions useful chrome driving face facebook

Cluster 1: world cup young champions crowned french men qatar fifa played


Predictions of new documents
Cluster of doc1 is  [0]
Cluster of doc2 is  [1]
```

## **PRACTICAL - 8**

**Aim:** Implement PoS Tagging on text.

### **Code:**

```
import nltk

text = "Mohamed Salah scored the fastest-ever Champions League hat-trick as Liverpool
turned it on in the second half to thrash Rangers."

from nltk.tokenize import word_tokenize

token_res = word_tokenize(text)

final_res = nltk.pos_tag(token_res)

print(final_res)
```

### **Output:**

```
[('Mohamed', 'NNP'),
 ('Salah', 'NNP'),
 ('scored', 'VBD'),
 ('the', 'DT'),
 ('fastest-ever', 'JJ'),
 ('Champions', 'NNP'),
 ('League', 'NNP'),
 ('hat-trick', 'NN'),
 ('as', 'IN'),
 ('Liverpool', 'NNP'),
 ('turned', 'VBD'),
 ('it', 'PRP'),
 ('on', 'IN'),
 ('in', 'IN'),
 ('the', 'DT'),
 ('second', 'JJ'),
 ('half', 'NN'),
 ('to', 'TO'),
 ('thrash', 'VB'),
 ('Rangers', 'NNP'),
 ('.', '.')] ]
```

## **PRACTICAL - 9**

**Aim:** Implement text processing with neural network

**Code:**

```
from keras_preprocessing.sequence import pad_sequences
from keras.layers import Embedding, LSTM, Dense, Dropout
from keras.preprocessing.text import Tokenizer
from keras.callbacks import EarlyStopping
from keras.models import Sequential
from tensorflow.keras.utils import to_categorical
import numpy as np

tokenizer = Tokenizer()

def dataset_preparation(data):
    corpus = data.lower().split("\n")
    tokenizer.fit_on_texts(corpus)
    total_words = len(tokenizer.word_index) + 1
    input_sequences = []
    for line in corpus:
        token_list = tokenizer.texts_to_sequences([line])[0]
        for i in range(1, len(token_list)):
            n_gram_sequence = token_list[:i+1]
            input_sequences.append(n_gram_sequence)
    max_sequence_len = max([len(x) for x in input_sequences])
    input_sequences = np.array(pad_sequences(input_sequences,
        maxlen=max_sequence_len, padding='pre'))
    predictors, label = input_sequences[:, :-1], input_sequences[:, -1]
    label = to_categorical(label, num_classes=total_words)
    return predictors, label, max_sequence_len, total_words
```



```
data = ""The cat and her kittens They put on their mittens, To eat a Christmas pie. The poor
little kittens. They lost their mittens, And then they began to cry. O mother dear, we sadly
fear We cannot go to-day, For we have lost our mittens. " "If it be so, ye shall not go, For ye
are naughty kittens.""
```

```
def create_model(predictors, label, max_sequence_len, total_words):
    model = Sequential()
    model.add(Embedding(total_words, 10, input_length=max_sequence_len-1))
    model.add(LSTM(150, return_sequences = True))
    model.add(LSTM(100))
    model.add(Dense(total_words, activation='softmax'))
    model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
    earlystop = EarlyStopping(monitor='val_loss', min_delta=0, patience=5, verbose=0,
mode='auto')
    model.fit(predictors, label, epochs=100, verbose=1, callbacks=[earlystop])
    print (model.summary())
    return model
```

```
import numpy as np
def generate_text(seed_text, next_words, max_sequence_len):
    for _ in range(next_words):
        token_list = tokenizer.texts_to_sequences([seed_text])[0]
        token_list = pad_sequences([token_list], maxlen=max_sequence_len-1,
padding='pre')
        predicted = np.argmax(model.predict(token_list, verbose=0))
        output_word = ""
        for word, index in tokenizer.word_index.items():
            if index == predicted:
                output_word = word
                break
        seed_text += " " + output_word
    return seed_text
predictors, label, max_sequence_len, total_words = dataset_preparation(data)
```

```
model = create_model(predictors, label, max_sequence_len, total_words)
text = generate_text("we have", 3, max_sequence_len)
print (text)
```

**Output:**

```
we have lost our mittens
```

## **PRACTICAL - 10**

**Aim:** Implement text processing with LSTM.

**Code:**

```

from google.colab import drive
drive.mount('/content/drive')

import pandas as pd
import numpy as np
from keras.utils.np_utils import to_categorical
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, Input
from keras.layers import Embedding, LSTM, Bidirectional, SimpleRNN, GRU
from keras.models import Sequential, Model

data = pd.read_csv('SMSSpamCollection', sep = '\t', names = ['label','message'])

text = data['message']
class_label = data['label']
classes_list = ["ham","spam"]
label_index = class_label.apply(classes_list.index)
label1 = np.asarray(label_index)
label = to_categorical(np.asarray(label1))

tk=Tokenizer(filters='!"#$%&()*+,-./:;<=>?@[\\]^_`{|}~\t\n',lower=True,split=" ")
tk.fit_on_texts(text)
index=tk.word_index
x = tk.texts_to_sequences(text)
vocab_size = len(index)

```

```
embedding_vecor_length = 32

padded_docs = pad_sequences(x, maxlen=embedding_vecor_length, padding='post')

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(padded_docs, label, test_size=0.30,
random_state=42)

lstm_input= Input(shape=(embedding_vecor_length,), dtype='int32', name='lstm_input')

x= Embedding(vocab_size+1, 100,
input_length=embedding_vecor_length,trainable=True)(lstm_input)

x1=LSTM(256,return_sequences=True)(x)

lstm_out= LSTM(128,return_sequences=False)(x1)

main_output = Dense(2,activation='softmax', name='main_output')(lstm_out)

model = Model(inputs=lstm_input, outputs=main_output)

model.compile(loss='binary_crossentropy', optimizer='Adam', metrics=['accuracy'])

print(model.summary())

model.fit(X_train, y_train, validation_data = (X_test, y_test),epochs=2,batch_size=100,
verbose=2)

predictions_test = model.predict(X_test)

predictions_test1 = np.zeros_like(predictions_test)

predictions_test1[np.arange(len(predictions_test)), predictions_test.argmax(1)] = 1

from sklearn.metrics import classification_report

print(classification_report(y_test,predictions_test1))
```

## Output:

Model: "model"

Layer (type)	Output Shape	Param #
lstm_input (InputLayer)	[(None, 32)]	0
embedding_2 (Embedding)	(None, 32, 100)	901000
lstm_4 (LSTM)	(None, 32, 256)	365568
lstm_5 (LSTM)	(None, 128)	197120
main_output (Dense)	(None, 2)	258

=====  
 Total params: 1,463,946  
 Trainable params: 1,463,946  
 Non-trainable params: 0

None

### Clasification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	1448
1	0.94	0.96	0.95	224
micro avg	0.99	0.99	0.99	1672
macro avg	0.97	0.98	0.97	1672
weighted avg	0.99	0.99	0.99	1672
samples avg	0.99	0.99	0.99	1672

## **PRACTICAL - 11**

**Aim:** Implement HMM/CRF on sequence tagging task.

**Code:**

```
import nltk

nltk.download('treebank')

tagged_sentences = nltk.corpus.treebank.tagged_sents()

def features(sentence, index):
    """ sentence: [w1, w2, ...], index: the index of the word """
    return {
        'word': sentence[index],
        'is_first': index == 0,
        'is_last': index == len(sentence) - 1,
        'is_capitalized': sentence[index][0].upper() == sentence[index][0],
        'is_all_caps': sentence[index].upper() == sentence[index],
        'is_all_lower': sentence[index].lower() == sentence[index],
        'prefix-1': sentence[index][0],
        'prefix-2': sentence[index][:2],
        'prefix-3': sentence[index][:3],
        'suffix-1': sentence[index][-1],
        'suffix-2': sentence[index][-2:],
        'suffix-3': sentence[index][-3:],
        'prev_word': " if index == 0 else sentence[index - 1],
        'next_word': " if index == len(sentence) - 1 else sentence[index + 1],
        'has_hyphen': '-' in sentence[index],
        'is_numeric': sentence[index].isdigit(),
        'capitals_inside': sentence[index][1:].lower() != sentence[index][1:]
    }
```

```
from nltk.tag.util import untag

cutoff = int(.75 * len(tagged_sentences))
training_sentences = tagged_sentences[:cutoff]
test_sentences = tagged_sentences[cutoff:]

def transform_to_dataset(tagged_sentences):
    X, y = [], []
    for tagged in tagged_sentences:
        X.append([features(untag(tagged), index) for index in range(len(tagged))])
        y.append([tag for _, tag in tagged])
    return X, y

X_train, y_train = transform_to_dataset(training_sentences)
X_test, y_test = transform_to_dataset(test_sentences)

from sklearn_crfsuite import CRF
model = CRF()

try:
    model.fit(X_train, y_train)
except AttributeError:
    pass

from sklearn_crfsuite import metrics
y_pred = model.predict(X_test)
print(metrics.flat_accuracy_score(y_test, y_pred))

sentence = ['I', 'am', 'Yash', '!']

def pos_tag(sentence):
    sentence_features = [features(sentence, index) for index in range(len(sentence))]
    return list(zip(sentence, model.predict([sentence_features])[0]))
```

```
print(pos_tag(sentence))
```

**Output:**

```
0.9602683593122289  
[('I', 'PRP'), ('am', 'VBP'), ('Yash', 'NNP'), ('!', '.'), ('.', '.')]
```